



# Homework #1


GitHub: [https://github.com/pballou/ECGR\\_4106/tree/master/Homework](https://github.com/pballou/ECGR_4106/tree/master/Homework)

1. Develop a multi-layer perceptron with three hidden layers (you pick the dimensions of the hidden layers) for the CIFAR-10 dataset.
  - a. Train the model from scratch (with randomized parameters) and plot the results (training loss and accuracy, validation accuracy) after 20 epochs.




- i. Does your network need more epochs for full training?  
**Yes, it is clearly still in the process of converging.**
- ii. Do you observe overfitting?  
**Not too much, as validation is changing with training.**
- iii. Make sure to save the trained parameters and model. Report and plot your training and validation results. Report precision, recall, F1 score, and confusion matrix.

	precision	recall	f1-score
0	0.53	0.45	0.48
1	0.46	0.54	0.49
2	0.34	0.22	0.27
3	0.30	0.23	0.26
4	0.37	0.32	0.34
5	0.39	0.34	0.36
6	0.38	0.54	0.45
7	0.42	0.46	0.44
8	0.47	0.58	0.52
9	0.47	0.49	0.48
accuracy			0.42
macro avg	0.41	0.42	0.41
weighted avg	0.41	0.42	0.41



- b. Explore the complexity of the network by increasing its width and depth.
- i. How do the training and validation results change?

**They both got worse. Fully connected networks are not very good with structured data like images.**



- ii. Compare them against the baseline. Do you see any overfitting?
- Not much, but the model is much worse overall.


2. Please implement the following steps for the housing dataset we overviewed during the lectures.
  - a. Build a multi-perceptron network that regresses the housing price (based on 20%, 80% split). Use the same number of features we did in the lecture without one-hot encoding. Please plot the training and validation results and report final accuracy and model complexity.

**I used a 3-layer network and here are the results:**

**Final Training Loss: 3.032738026819731**


**Final Validation Loss: 8.515102410316468**

**Final Validation RMSE: 0.3818462356939064**



- b. Build a multi-perceptron network that regresses the housing price (based on 20%, 80% split). Use the same number of features we did in the lecture, but this time also add one-hot encoding. Please plot the training and validation results and report the final accuracy and model complexity.

Final Training Loss: 0.03129004304738421  
 Final Validation Loss: 1.0817846432328224  
 Final Validation RMSE: 1.0400887662804106



- i. Do you see meaningful changes against 2.a?

**Yes, using one-hot encoding clearly had a positive impact on the results. Only the RMSE got a little worse, but both training and validation loss decreased significantly, showing that one-hot encoding helped the model learn the data better.**

- c. Increase the complexity of the network for problem 2.c. and compare your results against 2.b.

**I added a couple more layers and was able to get some better results.**

**Although the training loss got slightly worse, both validation loss and RMSE got better, so overall it performed more favorably.**

Final Training Loss: 0.14796865466786058  
 Final Validation Loss: 0.8522834584116936  
 Final Validation RMSE: 0.12080519385081238

