

Problem_1

February 7, 2024

```
[ ]: '''  
Patrick Ballou  
ID: 801130521  
ECGR 4106  
Homework 1  
Problem 1  
'''
```

```
[ ]: '\nPatrick Ballou\nID: 801130521\nECGR 4106\nHomework 1\nProblem 1\n'
```

```
[ ]: import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch import cuda  
from torchvision import datasets, transforms  
from torch.utils.data import DataLoader  
import numpy as np  
from sklearn import metrics  
import matplotlib.pyplot as plt  
import seaborn as sns  
from tqdm.notebook import tqdm  
import ssl  
ssl._create_default_https_context = ssl._create_unverified_context
```

```
[ ]: #check if GPU is available and set the device accordingly  
#device = 'torch.device("cuda:0" if torch.cuda.is_available() else "cpu")'  
device = 'cuda'  
print("Using GPU: ", cuda.get_device_name())  
  
gpu_info = !nvidia-smi  
gpu_info = '\n'.join(gpu_info)  
if gpu_info.find('failed') >= 0:  
    print('Not connected to a GPU')  
else:  
    print(gpu_info)
```

Using GPU: Quadro T2000
Wed Feb 7 02:18:31 2024

```

+-----+
-----+
| NVIDIA-SMI 551.23           Driver Version: 551.23           CUDA Version:
12.4           |
|-----+-----+-----+
-----+
| GPU Name                   TCC/WDDM | Bus-Id           Disp.A | Volatile
Uncorr. ECC |
| Fan Temp   Perf          Pwr:Usage/Cap |           Memory-Usage | GPU-Util
Compute M. |
|                   |                   |
MIG M. |
|=====+=====+=====
=====|
|  0  Quadro T2000           WDDM |  00000000:01:00.0  On |
N/A |
| N/A   41C    P8           6W /   60W |   1621MiB /   4096MiB |    8%
Default |
|                   |                   |
N/A |
+-----+-----+-----+
-----+

+-----+
-----+
| Processes:
|
| GPU   GI    CI          PID    Type    Process name
GPU Memory |
|       ID    ID
Usage      |
|=====+=====+=====
=====|
|  0    N/A   N/A       1520    C+G    ...les\Microsoft OneDrive\OneDrive.exe
N/A      |
|  0    N/A   N/A       4868    C+G    ...siveControlPanel\SystemSettings.exe
N/A      |
|  0    N/A   N/A       5960    C+G    ...b3d8bbwe\Microsoft.Media.Player.exe
N/A      |
|  0    N/A   N/A      14300    C+G    C:\Windows\explorer.exe
N/A      |
|  0    N/A   N/A      15240    C+G    ...Programs\Microsoft VS Code\Code.exe
N/A      |
|  0    N/A   N/A      15536    C+G    ...Brave-Browser\Application\brave.exe
N/A      |
|  0    N/A   N/A      15812    C+G    ...21.0_x64__8wekyb3d8bbwe\GameBar.exe
N/A      |
|  0    N/A   N/A      16124    C+G    ...Search_cw5n1h2txyewy\SearchApp.exe

```

N/A						
	0	N/A	N/A	17904	C+G	...t.LockApp_cw5n1h2txyewy\LockApp.exe
N/A						
	0	N/A	N/A	18012	C	...ri\anaconda3\envs\dl_env\python.exe
N/A						
	0	N/A	N/A	18112	C+G	...x64__rz1tebttb220\DolbyAccess.exe
N/A						
	0	N/A	N/A	18264	C+G	...AppData\Roaming\Spotify\Spotify.exe
N/A						
	0	N/A	N/A	18540	C+G	...ekyb3d8bbwe\PhoneExperienceHost.exe
N/A						
	0	N/A	N/A	19408	C+G	...CBS_cw5n1h2txyewy\TextInputHost.exe
N/A						
	0	N/A	N/A	19940	C+G	...aam7r\AcrobatNotificationClient.exe
N/A						
	0	N/A	N/A	20320	C+G	... Stream\85.0.37.0\GoogleDriveFS.exe
N/A						
	0	N/A	N/A	20376	C+G	...5n1h2txyewy\ShellExperienceHost.exe
N/A						
	0	N/A	N/A	23004	C	...ri\anaconda3\envs\dl_env\python.exe
N/A						
	0	N/A	N/A	24568	C+G	...ta\Local\Programs\Notion\Notion.exe
N/A						

```
+-----+
-----+
```

1 Problem 1A:

```
[ ]: data_path = '../.. /Datasets'
normalizing_set = datasets.CIFAR10(root=data_path, train=True, download=True,
↳ transform=transforms.ToTensor())
```

Files already downloaded and verified

```
[ ]: images = torch.stack([img_t for img_t, _ in normalizing_set], dim=3)
mean = images.view(3, -1).mean(dim=1)
std = images.view(3, -1).std(dim=1)

normalizer = transforms.Compose([transforms.ToTensor(), transforms.
↳ Normalize(mean,std)])

cifar_train = datasets.CIFAR10(root=data_path, train=True, transform=normalizer)
cifar_val = datasets.CIFAR10(root=data_path, train=False, transform=normalizer)

train_loader = DataLoader(cifar_train, batch_size=64, shuffle=True)
val_loader = DataLoader(cifar_val, batch_size=64, shuffle=False)
```

```
[ ]: class CifarNet(nn.Module):
    def __init__(self):
        super(CifarNet, self).__init__()
        self.flatten = nn.Flatten()
        self.relu = nn.ReLU()
        self.fc1 = nn.Linear(32*32*3, 512)
        self.fc2 = nn.Linear(512, 128)
        self.fc3 = nn.Linear(128, 32)
        self.fc4 = nn.Linear(32, 10)

    def forward(self, x):
        x = self.flatten(x)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.relu(self.fc3(x))
        x = self.fc4(x)
        return x

model = CifarNet().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001)

epochs = 20

training_losses = []
training_accuracies = []
validation_losses = []
validation_accuracies = []

[ ]: for epoch in range(epochs):
    model.train()

    running_loss = 0.0
    correct = 0
    total = 0

    with tqdm(total=len(train_loader), desc=f'Epoch {epoch + 1}/{epochs}',
        ↪unit=' batch') as pbar:
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
```

```

        pbar.update()

        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    training_losses.append(running_loss / len(train_loader))
    training_accuracies.append(100 * correct / total)

    model.eval()

    running_loss = 0.0
    correct = 0
    total = 0
    predicted_labels = []
    true_labels = []

    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)
            running_loss += loss.item()

            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            predicted_labels.extend(predicted.cpu().numpy())
            true_labels.extend(labels.cpu().numpy())

    validation_losses.append(running_loss / len(val_loader))
    validation_accuracies.append(100 * correct / total)

    pbar.set_postfix({'Training Loss ': training_losses[-1], 'Validation_
↪Loss ': validation_losses[-1]})

torch.save(model.state_dict(), '../..Models/hw1_1a.pth')

```

```

Epoch 1/20:  0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 2/20:  0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 3/20:  0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 4/20:  0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 5/20:  0%|          | 0/782 [00:00<?, ? batch/s]

```

```
Epoch 6/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 7/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 8/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 9/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 10/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 11/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 12/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 13/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 14/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 15/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 16/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 17/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 18/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 19/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 20/20: 0%|         | 0/782 [00:00<?, ? batch/s]
```

```
[ ]: print("Final Training Loss:", training_losses[-1])
      print("Final Validation Loss:", validation_losses[-1])
      print("Final Training Accuracy:", training_accuracies[-1])
      print("Final Validation Accuracy:", validation_accuracies[-1])
      epochs_range = range(1, epochs + 1)

      plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(epochs_range, training_losses, label='Training Loss')
      plt.plot(epochs_range, validation_losses, label='Validation Loss')
      plt.title('Training and Validation Loss')
      plt.xlabel('Epoch')
      plt.ylabel('Loss')
      plt.legend()

      plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 2)
      plt.plot(epochs_range, training_accuracies, label='Training Accuracy')
      plt.plot(epochs_range, validation_accuracies, label='Validation Accuracy')
      plt.title('Training and Validation Accuracy')
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
      plt.legend()
```

```

confusion_matrix = metrics.confusion_matrix(true_labels, predicted_labels)
print(metrics.classification_report(true_labels, predicted_labels))

class_labels = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
plt.figure(figsize=(12, 8))
sns.heatmap(confusion_matrix, annot=True, fmt='d', xticklabels=class_labels, yticklabels=class_labels)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```



Final Training Loss: 1.6622318390690152


Final Validation Loss: 1.6518301948620255

Final Training Accuracy: 41.076

Final Validation Accuracy: 41.71

	precision	recall	f1-score	support
0	0.53	0.45	0.48	1000
1	0.46	0.54	0.49	1000
2	0.34	0.22	0.27	1000
3	0.30	0.23	0.26	1000
4	0.37	0.32	0.34	1000
5	0.39	0.34	0.36	1000
6	0.38	0.54	0.45	1000
7	0.42	0.46	0.44	1000
8	0.47	0.58	0.52	1000
9	0.47	0.49	0.48	1000
accuracy			0.42	10000
macro avg	0.41	0.42	0.41	10000
weighted avg	0.41	0.42	0.41	10000





2 Problem 1B:

```
[ ]: data_path = '../..../Datasets'
normalizing_set = datasets.CIFAR10(root=data_path, train=True, download=True,
↳ transform=transforms.ToTensor())
```

Files already downloaded and verified

```
[ ]: images = torch.stack([img_t for img_t, _ in normalizing_set], dim=3)
mean = images.view(3, -1).mean(dim=1)
std = images.view(3, -1).std(dim=1)

normalizer = transforms.Compose([transforms.ToTensor(), transforms.
↳ Normalize(mean,std)])

cifar_train = datasets.CIFAR10(root=data_path, train=True, transform=normalizer)
cifar_val = datasets.CIFAR10(root=data_path, train=False, transform=normalizer)

train_loader = DataLoader(cifar_train, batch_size=64, shuffle=True)
val_loader = DataLoader(cifar_val, batch_size=64, shuffle=False)
```

```
[ ]: class CifarNet(nn.Module):
    def __init__(self):
        super(CifarNet, self).__init__()
        self.flatten = nn.Flatten()
        self.relu = nn.ReLU()
        self.fc1 = nn.Linear(32*32*3, 1024)
        self.fc2 = nn.Linear(1024, 256)
        self.fc3 = nn.Linear(256, 128)
        self.fc4 = nn.Linear(128, 64)
        self.fc5 = nn.Linear(64, 16)
        self.fc6 = nn.Linear(16, 10)

    def forward(self, x):
        x = self.flatten(x)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.relu(self.fc3(x))
        x = self.relu(self.fc4(x))
        x = self.relu(self.fc5(x))
        x = self.fc6(x)
        return x

model = CifarNet().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001)

epochs = 20

training_losses = []
training_accuracies = []
validation_losses = []
validation_accuracies = []

[ ]: for epoch in range(epochs):
    model.train()

    running_loss = 0.0
    correct = 0
    total = 0

    with tqdm(total=len(train_loader), desc=f'Epoch {epoch + 1}/{epochs}',
        ↪unit=' batch') as pbar:
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
```

```

        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
        pbar.update()

    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

training_losses.append(running_loss / len(train_loader))
training_accuracies.append(100 * correct / total)

model.eval()

running_loss = 0.0
correct = 0
total = 0
predicted_labels = []
true_labels = []

with torch.no_grad():
    for inputs, labels in val_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        outputs = model(inputs)
        loss = criterion(outputs, labels)
        running_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        predicted_labels.extend(predicted.cpu().numpy())
        true_labels.extend(labels.cpu().numpy())

validation_losses.append(running_loss / len(val_loader))
validation_accuracies.append(100 * correct / total)

pbar.set_postfix({'Training Loss ': training_losses[-1], 'Validation_
↪Loss ': validation_losses[-1]})

torch.save(model.state_dict(), '../Models/hw1_1b.pth')

```

```

Epoch 1/20:   0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 2/20:   0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 3/20:   0%|          | 0/782 [00:00<?, ? batch/s]

```

```
Epoch 4/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 5/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 6/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 7/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 8/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 9/20: 0%|          | 0/782 [00:00<?, ? batch/s]
Epoch 10/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 11/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 12/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 13/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 14/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 15/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 16/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 17/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 18/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 19/20: 0%|         | 0/782 [00:00<?, ? batch/s]
Epoch 20/20: 0%|         | 0/782 [00:00<?, ? batch/s]
```

```
[ ]: print("Final Training Loss:", training_losses[-1])
      print("Final Validation Loss:", validation_losses[-1])
      print("Final Training Accuracy:", training_accuracies[-1])
      print("Final Validation Accuracy:", validation_accuracies[-1])
      epochs_range = range(1, epochs + 1)

      plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(epochs_range, training_losses, label='Training Loss')
      plt.plot(epochs_range, validation_losses, label='Validation Loss')
      plt.title('Training and Validation Loss')
      plt.xlabel('Epoch')
      plt.ylabel('Loss')
      plt.legend()

      plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 2)
      plt.plot(epochs_range, training_accuracies, label='Training Accuracy')
      plt.plot(epochs_range, validation_accuracies, label='Validation Accuracy')
      plt.title('Training and Validation Accuracy')
      plt.xlabel('Epoch')
```

```

plt.ylabel('Accuracy')
plt.legend()

confusion_matrix = metrics.confusion_matrix(true_labels, predicted_labels)
print(metrics.classification_report(true_labels, predicted_labels))

class_labels = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
plt.figure(figsize=(12, 8))
sns.heatmap(confusion_matrix, annot=True, fmt='d', xticklabels=class_labels, yticklabels=class_labels)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()

```

Final Training Loss: 2.103359831720972

Final Validation Loss: 2.0812858092557094

Final Training Accuracy: 21.832

Final Validation Accuracy: 22.46

	precision	recall	f1-score	support
0	0.35	0.47	0.40	1000
1	0.28	0.01	0.02	1000
2	0.26	0.04	0.08	1000
3	0.15	0.50	0.23	1000
4	0.00	0.00	0.00	1000
5	0.12	0.16	0.14	1000
6	0.24	0.01	0.02	1000
7	0.67	0.00	0.00	1000
8	0.31	0.47	0.37	1000
9	0.26	0.58	0.36	1000
accuracy			0.22	10000
macro avg	0.26	0.22	0.16	10000
weighted avg	0.26	0.22	0.16	10000

```

c:\Users\patri\anaconda3\envs\dl_env\Lib\site-
packages\sklearn\metrics\_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\patri\anaconda3\envs\dl_env\Lib\site-
packages\sklearn\metrics\_classification.py:1497: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```
c:\Users\patri\anaconda3\envs\dl_env\Lib\site-  
packages\sklearn\metrics\_classification.py:1497: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

