

# Problem\_2

March 17, 2024

```
[ ]: '''  
Patrick Ballou  
ID: 801130521  
ECGR 4106  
Homework 3  
Problem 2  
'''
```

```
[ ]: '\nPatrick Ballou\nID: 801130521\nECGR 4106\nHomework 3\nProblem 2\n'
```

```
[ ]: import torch  
import torch.nn as nn  
import torch.optim as optim  
import time  
import requests  
from torch import cuda  
from torchvision import datasets, transforms  
from torch.utils.data import DataLoader, Dataset  
import numpy as np  
import pandas as pd  
from sklearn import metrics  
from sklearn.preprocessing import StandardScaler as SS  
from sklearn.model_selection import train_test_split
```

```
[ ]: #check if GPU is available and set the device accordingly  
#device = 'torch.device("cuda:0" if torch.cuda.is_available() else "cpu")'  
device = 'cuda'  
print("Using GPU: ", cuda.get_device_name())  
  
gpu_info = !nvidia-smi  
gpu_info = '\n'.join(gpu_info)  
if gpu_info.find('failed') >= 0:  
    print('Not connected to a GPU')  
else:  
    print(gpu_info)
```

Using GPU: Quadro T2000  
Sun Mar 17 19:44:57 2024

```

+-----+
-----+
| NVIDIA-SMI 551.23           Driver Version: 551.23           CUDA Version:
12.4           |
|-----+-----+-----+
-----+
| GPU Name                    TCC/WDDM | Bus-Id            Disp.A | Volatile
Uncorr. ECC |
| Fan Temp    Perf          Pwr:Usage/Cap |            Memory-Usage | GPU-Util
Compute M. |
|                    |                    |
MIG M. |
|=====+=====+=====
=====|
|  0  Quadro T2000          WDDM |  00000000:01:00.0  On |
N/A |
| N/A   55C    P8           5W /   30W |    575MiB /   4096MiB |    16%
Default |
|                    |                    |
N/A |
+-----+-----+-----+
-----+

+-----+
-----+
| Processes:
|
| GPU  GI  CI           PID  Type  Process name
GPU Memory |
|      ID  ID
Usage      |
|=====+=====+=====
=====|
|  0  N/A  N/A        1788   C+G   ...b3d8bbwe\Microsoft.Media.Player.exe
N/A      |
|  0  N/A  N/A        5084   C+G   ...siveControlPanel\SystemSettings.exe
N/A      |
|  0  N/A  N/A        7208   C+G   ...41.0_x64__8wekyb3d8bbwe\GameBar.exe
N/A      |
|  0  N/A  N/A       10352   C+G   ...e Stream\88.0.0.0\GoogleDriveFS.exe
N/A      |
|  0  N/A  N/A       12968   C+G   ...t.LockApp_cw5n1h2txyewy\LockApp.exe
N/A      |
|  0  N/A  N/A       14388   C+G   C:\Windows\explorer.exe
N/A      |
|  0  N/A  N/A       14988   C+G   ...ekyb3d8bbwe\PhoneExperienceHost.exe
N/A      |
|  0  N/A  N/A       15592   C+G   ...Programs\Microsoft VS Code\Code.exe

```

N/A						
	0	N/A	N/A	15932	C+G	...Search_cw5n1h2txyewy\SearchApp.exe
N/A						
	0	N/A	N/A	17808	C+G	...1300.0_x64__8j3eq9eme6ctt\IGCC.exe
N/A						
	0	N/A	N/A	18876	C+G	...CBS_cw5n1h2txyewy\TextInputHost.exe
N/A						
	0	N/A	N/A	19076	C+G	...ta\Local\Programs\Notion\Notion.exe
N/A						
	0	N/A	N/A	19888	C+G	...Brave-Browser\Application\brave.exe
N/A						
	0	N/A	N/A	20148	C+G	...aam7r\AcrobatNotificationClient.exe
N/A						
	0	N/A	N/A	21240	C+G	...les\Microsoft OneDrive\OneDrive.exe
N/A						
	0	N/A	N/A	21896	C+G	...5n1h2txyewy\ShellExperienceHost.exe
N/A						

+-----+  
-----+

```
[ ]: # Pred dataset for training
url = "https://raw.githubusercontent.com/karpathy/char-rnn/master/data/
      ↪tinyshakespeare/input.txt"
response = requests.get(url)
text = response.text # This is the entire text data

chars = sorted(list(set(text)))
char_to_int = {ch: i for i, ch in enumerate(chars)}
int_to_char = {i: ch for i, ch in enumerate(chars)}

# Encode the text into integers
encoded_text = [char_to_int[ch] for ch in text]
```

# 1 Problem 2A: LSTM(20, 30, 50)

```
[ ]: sequence_length = 20

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
```

```

sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)

# Dataset class
class CharDataset(Dataset):
    def __init__(self, sequences, targets):
        self.sequences = sequences
        self.targets = targets

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, index):
        return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset, [
    ↪train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the LSTM model
class CharLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharLSTM, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.lstm = nn.LSTM(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.lstm(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device),
                torch.zeros(1, batch_size, self.hidden_size, device=device))

```

```

# Hyperparameters
input_size = len(chars)
hidden_size = 256
output_size = len(chars)
learning_rate = 0.001

# Model, loss, and optimizer
model = CharLSTM(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

epochs = 25

init_time = time.time()
print("20 sequence LSTM results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

```

```
print(f"Accuracy of test set: {100 * correct / total}%")

#torch.save(model.state_dict(), '../..Models/hw3_2a_20.pth')
```

20 sequence LSTM results:

```
Epoch 1, Training Loss: 1.7039286458348653
Epoch 2, Training Loss: 1.4828566941152137
Epoch 3, Training Loss: 1.4228907150060912
Epoch 4, Training Loss: 1.3878118183188102
Epoch 5, Training Loss: 1.361920929226941
Epoch 6, Training Loss: 1.3428072792254024
Epoch 7, Training Loss: 1.3274340351217448
Epoch 8, Training Loss: 1.3143731060636092
Epoch 9, Training Loss: 1.303149068528253
Epoch 10, Training Loss: 1.2939258513953997
Epoch 11, Training Loss: 1.2853256974986964
Epoch 12, Training Loss: 1.2785495665657471
Epoch 13, Training Loss: 1.2721631698955835
Epoch 14, Training Loss: 1.2667568558447675
Epoch 15, Training Loss: 1.2605822635873323
Epoch 16, Training Loss: 1.2563257013852065
Epoch 17, Training Loss: 1.2524168876417054
Epoch 18, Training Loss: 1.2479410454827482
Epoch 19, Training Loss: 1.2448831394066842
Epoch 20, Training Loss: 1.2421532968594573
Epoch 21, Training Loss: 1.2385113516280655
Epoch 22, Training Loss: 1.2367871534068626
Epoch 23, Training Loss: 1.2343798212224015
Epoch 24, Training Loss: 1.232003951672843
Epoch 25, Training Loss: 1.2305857209350894
Training time: 10.025756839911143 minutes
Accuracy of test set: 58.42608988008517%
```

```
[ ]: sequence_length = 30

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)
```

```

# Dataset class
class CharDataset(Dataset):
    def __init__(self, sequences, targets):
        self.sequences = sequences
        self.targets = targets

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, index):
        return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset,
    ↪ [train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the LSTM model
class CharLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharLSTM, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.lstm = nn.LSTM(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.lstm(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device),
            torch.zeros(1, batch_size, self.hidden_size, device=device))

# Hyperparameters
input_size = len(chars)

```

```

hidden_size = 256
output_size = len(chars)
learning_rate = 0.001

# Model, loss, and optimizer
model = CharLSTM(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

epochs = 25

init_time = time.time()
print("30 sequence LSTM results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

print(f"Accuracy of test set: {100 * correct / total}%")

```



```
#torch.save(model.state_dict(), '../..Models/hw3_2a_30.pth')
```

30 sequence LSTM results:

Epoch 1, Training Loss: 1.6993088538251386  
Epoch 2, Training Loss: 1.474142562501812  
Epoch 3, Training Loss: 1.4137238254679803  
Epoch 4, Training Loss: 1.3783982019963377  
Epoch 5, Training Loss: 1.3537912187788996  
Epoch 6, Training Loss: 1.3340486319531535  
Epoch 7, Training Loss: 1.3190894767094876  
Epoch 8, Training Loss: 1.3062398333403822  
Epoch 9, Training Loss: 1.2956803129516794  
Epoch 10, Training Loss: 1.2862978480538279  
Epoch 11, Training Loss: 1.278085498514356  
Epoch 12, Training Loss: 1.270350139226916  
Epoch 13, Training Loss: 1.263484604259161  
Epoch 14, Training Loss: 1.2579474826260537  
Epoch 15, Training Loss: 1.252641338822969  
Epoch 16, Training Loss: 1.2477584728480617  
Epoch 17, Training Loss: 1.2438362084085477  
Epoch 18, Training Loss: 1.2396418672574367  
Epoch 19, Training Loss: 1.235707599415358  
Epoch 20, Training Loss: 1.2322622239076741  
Epoch 21, Training Loss: 1.2298383069188856  
Epoch 22, Training Loss: 1.2269100274525184  
Epoch 23, Training Loss: 1.225325421208427  
Epoch 24, Training Loss: 1.2233063099255044  
Epoch 25, Training Loss: 1.2213229424075975  
Training time: 11.263746031125386 minutes  
Accuracy of test set: 58.58575443912979%

```
[ ]: sequence_length = 50

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)

# Dataset class
```

```

class CharDataset(Dataset):
    def __init__(self, sequences, targets):
        self.sequences = sequences
        self.targets = targets

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, index):
        return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset, [
    ↪ train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the LSTM model
class CharLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharLSTM, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.lstm = nn.LSTM(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.lstm(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device),
                torch.zeros(1, batch_size, self.hidden_size, device=device))

# Hyperparameters
input_size = len(chars)
hidden_size = 256
output_size = len(chars)

```

```

learning_rate = 0.001

# Model, loss, and optimizer
model = CharLSTM(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

epochs = 25

init_time = time.time()
print("50 sequence LSTM results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

print(f"Accuracy of test set: {100 * correct / total}%")

#torch.save(model.state_dict(), '../..Models/hw3_2a_50.pth')

```

50 sequence LSTM results:

Epoch 1, Training Loss: 1.7033157834293955  
 Epoch 2, Training Loss: 1.472816352770021  
 Epoch 3, Training Loss: 1.4091637182092072  
 Epoch 4, Training Loss: 1.3726160892830093  
 Epoch 5, Training Loss: 1.346235753165959  
 Epoch 6, Training Loss: 1.325980530674737  
 Epoch 7, Training Loss: 1.3088446769271789  
 Epoch 8, Training Loss: 1.2956975393650298  
 Epoch 9, Training Loss: 1.284235503269634  
 Epoch 10, Training Loss: 1.2747833096490258  
 Epoch 11, Training Loss: 1.2659264016383998  
 Epoch 12, Training Loss: 1.2586697097950559  
 Epoch 13, Training Loss: 1.2507322769119893  
 Epoch 14, Training Loss: 1.245463993572946  
 Epoch 15, Training Loss: 1.240341373383272  
 Epoch 16, Training Loss: 1.234599030893411  
 Epoch 17, Training Loss: 1.2303823452566331  
 Epoch 18, Training Loss: 1.227314035949945  
 Epoch 19, Training Loss: 1.2234475778517602  
 Epoch 20, Training Loss: 1.220022144413049  
 Epoch 21, Training Loss: 1.217116185521279  
 Epoch 22, Training Loss: 1.21419994343885  
 Epoch 23, Training Loss: 1.212461400247513  
 Epoch 24, Training Loss: 1.2099621501052384  
 Epoch 25, Training Loss: 1.2077432776297752  
 Training time: 12.90316569407781 minutes  
 Accuracy of test set: 59.213965185660044%

## 2 Problem 2B: GRU(20, 30, 50)

```
[ ]: sequence_length = 20

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)

# Dataset class
class CharDataset(Dataset):
```

```

def __init__(self, sequences, targets):
    self.sequences = sequences
    self.targets = targets

def __len__(self):
    return len(self.sequences)

def __getitem__(self, index):
    return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset, [
    ↪ train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the GRU model
class CharGRU(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharGRU, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.gru(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device))

# Hyperparameters
input_size = len(chars)
hidden_size = 256
output_size = len(chars)
learning_rate = 0.001

```

```

# Model, loss, and optimizer
model = CharGRU(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

epochs = 25

init_time = time.time()
print("20 sequence GRU results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

print(f"Accuracy of test set: {100 * correct / total}%")

#torch.save(model.state_dict(), '../Models/hw3_2b_20.pth')

```

20 sequence GRU results:

Epoch 1, Training Loss: 1.6904839627782435

Epoch 2, Training Loss: 1.497342645540637

Epoch 3, Training Loss: 1.4507262527993952  
Epoch 4, Training Loss: 1.4255671261270635  
Epoch 5, Training Loss: 1.4101337558807345  
Epoch 6, Training Loss: 1.3995701468373878  
Epoch 7, Training Loss: 1.3912026316047195  
Epoch 8, Training Loss: 1.3853752609391494  
Epoch 9, Training Loss: 1.3809091519399819  
Epoch 10, Training Loss: 1.376863594574622  
Epoch 11, Training Loss: 1.374466078185994  
Epoch 12, Training Loss: 1.3727443141288889  
Epoch 13, Training Loss: 1.3720990438794378  
Epoch 14, Training Loss: 1.3719647168873783  
Epoch 15, Training Loss: 1.3736919506110477  
Epoch 16, Training Loss: 1.3745446196675506  
Epoch 17, Training Loss: 1.3749985953411423  
Epoch 18, Training Loss: 1.3745096339723946  
Epoch 19, Training Loss: 1.3774950651348288  
Epoch 20, Training Loss: 1.3789241221090602  
Epoch 21, Training Loss: 1.3827220115573393  
Epoch 22, Training Loss: 1.3876126805342823  
Epoch 23, Training Loss: 1.3884858456720788  
Epoch 24, Training Loss: 1.3935831162396752  
Epoch 25, Training Loss: 1.3947730505100557  
Training time: 9.919338961442312 minutes  
Accuracy of test set: 55.9645859016026%

```
[ ]: sequence_length = 30

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)

# Dataset class
class CharDataset(Dataset):
    def __init__(self, sequences, targets):
        self.sequences = sequences
        self.targets = targets
```

```

def __len__(self):
    return len(self.sequences)

def __getitem__(self, index):
    return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset,
    ↪[train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the GRU model
class CharGRU(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharGRU, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.gru(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device))

# Hyperparameters
input_size = len(chars)
hidden_size = 256
output_size = len(chars)
learning_rate = 0.001

# Model, loss, and optimizer
model = CharGRU(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

```



```

epochs = 25

init_time = time.time()
print("30 sequence GRU results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

print(f"Accuracy of test set: {100 * correct / total}%")

#torch.save(model.state_dict(), '../..Models/hw3_2b_30.pth')

```

30 sequence GRU results:

```

Epoch 1, Training Loss: 1.688761063437317
Epoch 2, Training Loss: 1.4947071062188277
Epoch 3, Training Loss: 1.446230951023389
Epoch 4, Training Loss: 1.4202568271740783
Epoch 5, Training Loss: 1.403517935940129
Epoch 6, Training Loss: 1.3898934693649043

```

Epoch 7, Training Loss: 1.3804380048538718  
Epoch 8, Training Loss: 1.3759228813104498  
Epoch 9, Training Loss: 1.369669947578725  
Epoch 10, Training Loss: 1.367014287098413  
Epoch 11, Training Loss: 1.3641738207156946  
Epoch 12, Training Loss: 1.3612514548067791  
Epoch 13, Training Loss: 1.3606549097300944  
Epoch 14, Training Loss: 1.3582993330229853  
Epoch 15, Training Loss: 1.3583210253093136  
Epoch 16, Training Loss: 1.3583726284665345  
Epoch 17, Training Loss: 1.3611586110060474  
Epoch 18, Training Loss: 1.3605015400115335  
Epoch 19, Training Loss: 1.3630748972835312  
Epoch 20, Training Loss: 1.3640652118369347  
Epoch 21, Training Loss: 1.3667174100756165  
Epoch 22, Training Loss: 1.3682563400993522  
Epoch 23, Training Loss: 1.3699411446174123  
Epoch 24, Training Loss: 1.373385976365224  
Epoch 25, Training Loss: 1.3744303944445313  
Training time: 11.177410554885864 minutes  
Accuracy of test set: 56.34747369695122%

```
[ ]: sequence_length = 50

# Create sequences and targets
sequences = []
targets = []
for i in range(0, len(encoded_text) - sequence_length):
    seq = encoded_text[i:i+sequence_length]
    target = encoded_text[i+sequence_length]
    sequences.append(seq)
    targets.append(target)

# Convert lists to PyTorch tensors
sequences = torch.tensor(sequences, dtype=torch.long)
targets = torch.tensor(targets, dtype=torch.long)

# Dataset class
class CharDataset(Dataset):
    def __init__(self, sequences, targets):
        self.sequences = sequences
        self.targets = targets

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, index):
```

```

        return self.sequences[index], self.targets[index]

# Instantiate the dataset
dataset = CharDataset(sequences, targets)

# Create data loaders
batch_size = 128
train_size = int(len(dataset) * 0.8)
test_size = len(dataset) - train_size
train_dataset, test_dataset = torch.utils.data.random_split(dataset,
    ↪ [train_size, test_size])

train_loader = DataLoader(train_dataset, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_dataset, shuffle=False, batch_size=batch_size)

```

```

[ ]: # Defining the GRU model
class CharGRU(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(CharGRU, self).__init__()
        self.hidden_size = hidden_size
        self.embedding = nn.Embedding(input_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x, hidden):
        embedded = self.embedding(x)
        output, hidden = self.gru(embedded, hidden)
        output = self.fc(output[:, -1, :])
        return output, hidden

    def init_hidden(self, batch_size):
        return (torch.zeros(1, batch_size, self.hidden_size, device=device))

# Hyperparameters
input_size = len(chars)
hidden_size = 256
output_size = len(chars)
learning_rate = 0.001

# Model, loss, and optimizer
model = CharGRU(input_size, hidden_size, output_size).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

epochs = 25

init_time = time.time()

```

```

print("50 sequence GRU results:")

# Training the model
for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    print(f"Epoch {epoch+1}, Training Loss: {running_loss / len(train_loader)}")

print(f"Training time: {(time.time() - init_time)/60} minutes")

# Validation
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        hidden = model.init_hidden(data.size(0))
        output, hidden = model(data, hidden)
        _, predicted = torch.max(output.data, 1)
        total += target.size(0)
        correct += (predicted == target).sum().item()

print(f"Accuracy of test set: {100 * correct / total}%")

#torch.save(model.state_dict(), '../..Models/hw3_2b_50.pth')

```

50 sequence GRU results:

```

Epoch 1, Training Loss: 1.6804511425160316
Epoch 2, Training Loss: 1.4817302748678474
Epoch 3, Training Loss: 1.435106966865232
Epoch 4, Training Loss: 1.4087642558830709
Epoch 5, Training Loss: 1.3914834176548379
Epoch 6, Training Loss: 1.3802862838049597
Epoch 7, Training Loss: 1.3710225728881391
Epoch 8, Training Loss: 1.3649185802036712
Epoch 9, Training Loss: 1.3613395527092838
Epoch 10, Training Loss: 1.356445742951389

```

Epoch 11, Training Loss: 1.353771461371352  
Epoch 12, Training Loss: 1.353007037043486  
Epoch 13, Training Loss: 1.3511901474925885  
Epoch 14, Training Loss: 1.3503325411740206  
Epoch 15, Training Loss: 1.349965271555083  
Epoch 16, Training Loss: 1.3516816538403869  
Epoch 17, Training Loss: 1.351813011686902  
Epoch 18, Training Loss: 1.3515537315419819  
Epoch 19, Training Loss: 1.3527359983668301  
Epoch 20, Training Loss: 1.3570245269842562  
Epoch 21, Training Loss: 1.3570453758784233  
Epoch 22, Training Loss: 1.3602206598924027  
Epoch 23, Training Loss: 1.3637807675848526  
Epoch 24, Training Loss: 1.365017493560177  
Epoch 25, Training Loss: 1.368766647631153  
Training time: 12.810324263572692 minutes  
Accuracy of test set: 56.79632759370419%