

# Problem\_1

April 27, 2024

```
[ ]: '''  
Patrick Ballou  
ID: 801130521  
ECGR 4106  
Homework 5  
Problem 1  
'''
```

```
[ ]: '\nPatrick Ballou\nID: 801130521\nECGR 4106\nHomework 5\nProblem 1\n'
```

```
[ ]: import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch import cuda  
import matplotlib.pyplot as plt  
from torch.utils.data import DataLoader, Dataset  
from sklearn.model_selection import train_test_split  
import numpy as np  
import time
```

```
[ ]: #check if GPU is available and set the device accordingly  
#device = 'torch.device("cuda:0" if torch.cuda.is_available() else "cpu")'  
device = 'cuda'  
print("Using GPU: ", cuda.get_device_name())  
  
gpu_info = !nvidia-smi  
gpu_info = '\n'.join(gpu_info)  
if gpu_info.find('failed') >= 0:  
    print('Not connected to a GPU')  
else:  
    print(gpu_info)
```

Using GPU: NVIDIA L4

Thu Apr 25 14:48:27 2024

```
+-----+  
-----+  
| NVIDIA-SMI 535.104.05                Driver Version: 535.104.05    CUDA Version:  
12.2          |
```

```

|-----+-----+-----+
-----+
| GPU Name Persistence-M | Bus-Id Disp.A | Volatile
Uncorr. ECC |
| Fan Temp Perf Pwr:Usage/Cap | Memory-Usage | GPU-Util
Compute M. |
| | | |
MIG M. |
|=====+=====+=====
=====|
| 0 NVIDIA L4 Off | 00000000:00:03.0 Off |
0 |
| N/A 52C P8 13W / 72W | 4MiB / 23034MiB | 0%
Default |
| | |
N/A |
+-----+-----+-----+
-----+

+-----+
-----+
| Processes:
|
| GPU GI CI PID Type Process name GPU
Memory |
| ID ID
Usage |
|=====+=====+=====
=====|
| No running processes found
|
+-----+-----+-----+
-----+

```

```
[ ]: text = ""Next character prediction is a fundamental task in the field of
↳natural language processing (NLP) that involves predicting the next
↳character in a sequence of text based on the characters that precede it.
↳This task is essential for various applications, including text
↳auto-completion, spell checking, and even in the development of
↳sophisticated AI models capable of generating human-like text.
```

At its core, next character prediction relies on statistical models or deep  
↳learning algorithms to analyze a given sequence of text and predict which  
↳character is most likely to follow. These predictions are based on patterns  
↳and relationships learned from large datasets of text during the training  
↳phase of the model.

One of the most popular approaches to next character prediction involves the use of Recurrent Neural Networks (RNNs), and more specifically, a variant called Long Short-Term Memory (LSTM) networks. RNNs are particularly well-suited for sequential data like text, as they can maintain information in 'memory' about previous characters to inform the prediction of the next character. LSTM networks enhance this capability by being able to remember long-term dependencies, making them even more effective for next character prediction tasks.

Training a model for next character prediction involves feeding it large amounts of text data, allowing it to learn the probability of each character's appearance following a sequence of characters. During this training process, the model adjusts its parameters to minimize the difference between its predictions and the actual outcomes, thus improving its predictive accuracy over time.

Once trained, the model can be used to predict the next character in a given piece of text by considering the sequence of characters that precede it. This can enhance user experience in text editing software, improve efficiency in coding environments with auto-completion features, and enable more natural interactions with AI-based chatbots and virtual assistants.

In summary, next character prediction plays a crucial role in enhancing the capabilities of various NLP applications, making text-based interactions more efficient, accurate, and human-like. Through the use of advanced machine learning models like RNNs and LSTMs, next character prediction continues to evolve, opening new possibilities for the future of text-based technology. """

```
[ ]: # Creating character vocabulary
chars = sorted(list(set(text)))
ix_to_char = {i: ch for i, ch in enumerate(chars)}
char_to_ix = {ch: i for i, ch in enumerate(chars)}

[ ]: def prepare_dataset(max_length):
    X = []
    y = []
    for i in range(len(text) - max_length):
        sequence = text[i:i + max_length]
        label = text[i + max_length]
        X.append([char_to_ix[char] for char in sequence])
        y.append(char_to_ix[label])

    X = np.array(X)
    y = np.array(y)
```

*# Splitting the dataset into training and validation sets*

```
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

```
# Converting data to PyTorch tensors
X_train = torch.tensor(X_train, dtype=torch.long)
y_train = torch.tensor(y_train, dtype=torch.long)
X_val = torch.tensor(X_val, dtype=torch.long)
y_val = torch.tensor(y_val, dtype=torch.long)

return X_train, y_train, X_val, y_val
```

```
[ ]: # Defining the Transformer model
class CharTransformer(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, nhead):
        super(CharTransformer, self).__init__()
        self.embedding = nn.Embedding(input_size, hidden_size)
        encoder_layers = nn.TransformerEncoderLayer(hidden_size, nhead,
↪batch_first=True)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers,
↪num_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        embedded = self.embedding(x)
        transformer_output = self.transformer_encoder(embedded)
        output = self.fc(transformer_output[:, -1, :]) # Get the output of the
↪last Transformer block
        return output
```

```
[ ]: def training_loop(model, criterion, optimizer, epochs, X_train, y_train, X_val,
↪y_val, max_length):

    train_history = []
    val_history = []
    init_time = time.time()
    print(f"{max_length} sequence transformer results:")

    for epoch in range(epochs):
        model.train()
        optimizer.zero_grad()
        output = model(X_train)
        loss = criterion(output, y_train)
        loss.backward()
        optimizer.step()

        # Validation
        model.eval()
```

```

        with torch.no_grad():
            val_output = model(X_val)
            val_loss = criterion(val_output, y_val)
            _, predicted = torch.max(val_output, 1)
            val_accuracy = (predicted == y_val).float().mean()

        train_history.append(loss.item())
        val_history.append(val_accuracy.item())

        if (epoch+1) % 10 == 0:
            print(f'Epoch {epoch+1}, Loss: {train_history[-1]}, Validation Loss:
↪ {val_loss.item()}, Validation Accuracy: {val_history[-1]}')

        print(f"Training time: {time.time() - init_time} seconds")
        save_path = f'../Models/hw5_1_{max_length}.pth'
        #torch.save(model.state_dict(), save_path)

    return train_history, val_history

# Prediction function
def predict_next_char(model, initial_str, max_length):
    model.eval()
    with torch.no_grad():
        initial_input = torch.tensor([char_to_ix[c] for c in
↪ initial_str[-max_length:]], dtype=torch.long).unsqueeze(0)
        prediction = model(initial_input)
        predicted_index = torch.argmax(prediction, dim=1).item()
        return ix_to_char[int(predicted_index)]

def next_char_test(model, max_length):
    test_str = "This is a simple example to demonstrate how to predict the next
↪ char"
    predicted_char = predict_next_char(model, test_str, max_length)
    print(f"Predicted next character: '{predicted_char}'")

```

```

[ ]: # Hyperparameters
hidden_size = 256
num_layers = 4
nhead = 4

max_lengths = [10, 20, 30]
learning_rate = .0005
epochs = 100

```

```

[ ]: train_histories = []
val_histories = []

```

```

for window_size in max_lengths:
    model = CharTransformer(len(chars), hidden_size, len(chars), num_layers,
                             ↪nhead)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    X_train, y_train, X_val, y_val = prepare_dataset(window_size)

    train_history, val_history = training_loop(model, criterion, optimizer,
    ↪epochs, X_train, y_train, X_val, y_val, window_size)
    train_histories.append(train_history)
    val_histories.append(val_history)
    print("\n")

    next_char_test(model, window_size)
    print("\n\n") if window_size != 30 else None

```

10 sequence transformer results:

Epoch 10, Loss: 2.4080586433410645, Validation Loss: 2.5051023960113525,  
Validation Accuracy: 0.2704402506351471  
Epoch 20, Loss: 2.0444071292877197, Validation Loss: 2.4036715030670166,  
Validation Accuracy: 0.33752620220184326  
Epoch 30, Loss: 1.6987130641937256, Validation Loss: 2.373068332672119,  
Validation Accuracy: 0.3668763041496277  
Epoch 40, Loss: 1.3319175243377686, Validation Loss: 2.4580023288726807,  
Validation Accuracy: 0.3731656074523926  
Epoch 50, Loss: 0.9410110116004944, Validation Loss: 2.65513277053833,  
Validation Accuracy: 0.36477985978126526  
Epoch 60, Loss: 0.6164496541023254, Validation Loss: 2.9696104526519775,  
Validation Accuracy: 0.3542976975440979  
Epoch 70, Loss: 0.4114442467689514, Validation Loss: 3.1755635738372803,  
Validation Accuracy: 0.3731656074523926  
Epoch 80, Loss: 0.22664591670036316, Validation Loss: 3.445261240005493,  
Validation Accuracy: 0.3689727485179901  
Epoch 90, Loss: 0.17504791915416718, Validation Loss: 3.696992874145508,  
Validation Accuracy: 0.36477985978126526  
Epoch 100, Loss: 0.13712351024150848, Validation Loss: 3.8674631118774414,  
Validation Accuracy: 0.36477985978126526  
Training time: 182.34048581123352 seconds

Predicted next character: 'a'

20 sequence transformer results:

Epoch 10, Loss: 2.6871519088745117, Validation Loss: 2.7132623195648193,

Validation Accuracy: 0.21684210002422333  
Epoch 20, Loss: 2.321113109588623, Validation Loss: 2.5463039875030518,  
Validation Accuracy: 0.25473684072494507  
Epoch 30, Loss: 2.103867292404175, Validation Loss: 2.503826141357422,  
Validation Accuracy: 0.269473671913147  
Epoch 40, Loss: 1.8689113855361938, Validation Loss: 2.532803773880005,  
Validation Accuracy: 0.2757894694805145  
Epoch 50, Loss: 1.5970028638839722, Validation Loss: 2.6253719329833984,  
Validation Accuracy: 0.2926315665245056  
Epoch 60, Loss: 1.316619873046875, Validation Loss: 2.818438768386841,  
Validation Accuracy: 0.28421053290367126  
Epoch 70, Loss: 0.9813405871391296, Validation Loss: 3.059547185897827,  
Validation Accuracy: 0.2947368323802948  
Epoch 80, Loss: 0.6918473243713379, Validation Loss: 3.256744861602783,  
Validation Accuracy: 0.2800000011920929  
Epoch 90, Loss: 0.4566172659397125, Validation Loss: 3.5641286373138428,  
Validation Accuracy: 0.2905263304710388  
Epoch 100, Loss: 0.3032119870185852, Validation Loss: 3.799943685531616,  
Validation Accuracy: 0.2926315665245056  
Training time: 419.69847226142883 seconds

Predicted next character: 'a'

30 sequence transformer results:

Epoch 10, Loss: 2.540902614593506, Validation Loss: 2.7158145904541016,  
Validation Accuracy: 0.23467230796813965  
Epoch 20, Loss: 2.2777161598205566, Validation Loss: 2.582134962081909,  
Validation Accuracy: 0.22832980751991272  
Epoch 30, Loss: 2.111842155456543, Validation Loss: 2.5371346473693848,  
Validation Accuracy: 0.25581395626068115  
Epoch 40, Loss: 1.952932596206665, Validation Loss: 2.6494176387786865,  
Validation Accuracy: 0.2473572939634323  
Epoch 50, Loss: 1.7562140226364136, Validation Loss: 2.642242431640625,  
Validation Accuracy: 0.268498957157135  
Epoch 60, Loss: 1.4736355543136597, Validation Loss: 2.824350118637085,  
Validation Accuracy: 0.25792813301086426  
Epoch 70, Loss: 1.1992594003677368, Validation Loss: 3.067345142364502,  
Validation Accuracy: 0.25792813301086426  
Epoch 80, Loss: 0.9267798066139221, Validation Loss: 3.3158175945281982,  
Validation Accuracy: 0.24947145581245422  
Epoch 90, Loss: 0.6758770942687988, Validation Loss: 3.6170268058776855,  
Validation Accuracy: 0.2452431321144104  
Epoch 100, Loss: 0.4670465886592865, Validation Loss: 3.9806814193725586,  
Validation Accuracy: 0.25369977951049805  
Training time: 665.9915945529938 seconds

Predicted next character: 'e'

```
[ ]: # Plotting training losses
plt.figure(figsize=(10, 5))
for i, max_length in enumerate(max_lengths):
    plt.plot(train_histories[i], label=f"Context Window: {max_length}")
plt.xlabel("Epochs")
plt.ylabel("Training Loss")
plt.title("Training Losses over Epochs")
plt.legend()
plt.show()

# Plotting validation accuracies
plt.figure(figsize=(10, 5))
for i, max_length in enumerate(max_lengths):
    plt.plot(val_histories[i], label=f"Context Window: {max_length}")
plt.xlabel("Epochs")
plt.ylabel("Validation Accuracy")
plt.title("Validation Accuracies over Epochs")
plt.legend()
plt.show()
```





