



UNC CHARLOTTE

The WILLIAM STATES LEE COLLEGE *of* ENGINEERING

Introduction to ML

Lecture 7: Recurrent Neural Network

Hamed Tabkhi

Department of Electrical and Computer Engineering,
University of North Carolina Charlotte (UNCC)


htabkhiv@uncc.edu



UNC CHARLOTTE

Intro to RNNs

- Countless learning tasks require dealing with sequential data.
- Video captioning, speech synthesis, and music generation all require that models produce outputs consisting of sequences.
- In other domains, such as time series prediction, video analysis, and musical information retrieval, a model must learn from inputs that are sequences.
- Recurrent neural networks (RNNs) are deep learning models that capture the dynamics of sequences via *recurrent* connections, which can be thought of as cycles in the network of nodes.
- RNNs can be thought of as feedforward neural networks where each layer's parameters (both conventional and recurrent) are shared across



- On the left recurrent connections are depicted via cyclic edges. On the right, we unfold the RNN over time steps.

Intro to RNNs

State at any time step t (h_t) could be computed based on both the current input x_t and the previous state h_{t-1}

$$h_t = f(x_t, h_{t-1}).$$


RNNs rose to prominence as the default models for handling complex sequential structure in deep learning and remain staple models for sequential modeling to this day.


RNN Key insights:

While the inputs and targets for many fundamental tasks in machine learning cannot easily be represented as fixed length vectors, they can often nevertheless be represented as varying-length sequences of fixed length vectors.

For example, videos can be represented as varying-length sequences of still images.

RNN presented





A, B and C are the parameters

RNN Formulation with Hidden States (Layers)

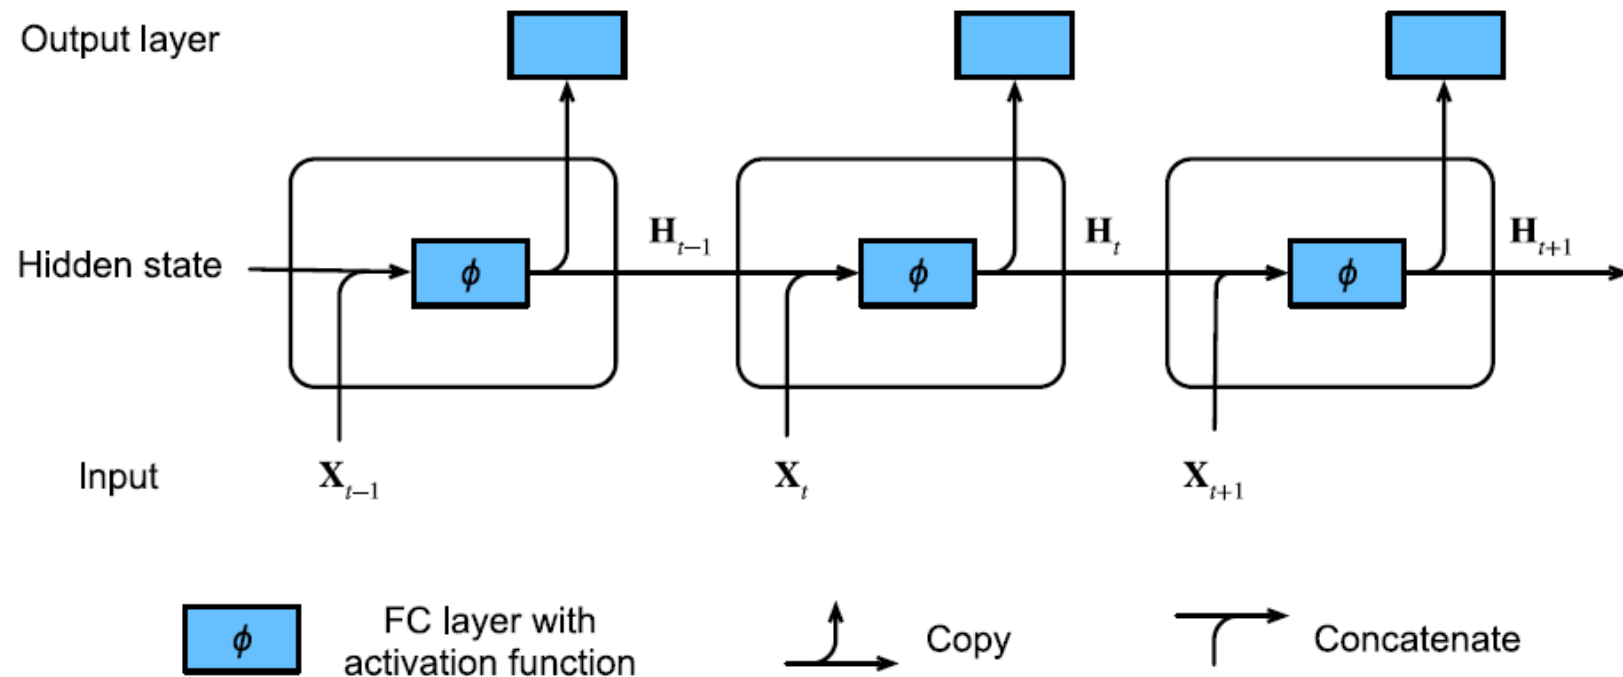


Fig illustrates the computational logic of an RNN at three adjacent time steps.

At any time step t , the computation of the hidden state can be treated as:

- (i) concatenating the input \mathbf{X}_t at the current time step t and the hidden state \mathbf{H}_{t-1} at the previous time step $t-1$
- (ii) feeding the concatenation result into a fully connected layer with the activation function ϕ .


The output of such a fully connected layer is the hidden state \mathbf{H}_t of the current time step t .

In this case, the model parameters are the concatenation of \mathbf{WXh} and \mathbf{Whh} , and a bias of \mathbf{Bh} .

What is more, \mathbf{H}_t will also be fed into the fully connected output layer to compute the output \mathbf{O}_t of the current time step t .

The hidden state of the current time step t , \mathbf{H}_t will participate in computing the hidden state \mathbf{H}_{t+1} of the next time step $t+1$.

Example: RNN for character-level language model



Predicting the next character based on the current and previous characters via an RNN for character-level language modeling

- Due to the recurrent computation of the hidden state in the
- hidden layer, the output of time step 3, is determined by the text sequence “m”, “a”, and “c”.
- Since the next character of the sequence in the training data is “h”, the correct prediction of time step 3 will depend on the probability distribution of the next character generated based on the feature sequence “m”, “a”, “c” and the target “h” of this time step.

